


# Exploring The Data Science Process



A conceptual diagram for Data Science. At the center is a white cloud containing the text "Data Science" in bold red font. Radiating from this central cloud are twelve lines, each ending in a white icon. The icons represent various data science concepts: a laptop with a bar chart and magnifying glass, a database cylinder, a cloud with up/down arrows, a document with a bar chart, a server rack, a line graph with a magnifying glass, a magnifying glass over a bar chart, three server units, a pie chart, a database cylinder with a padlock, a line graph with an upward arrow, and a database cylinder with a magnifying glass. The background is a solid reddish-pink color with faint white cloud shapes.

# Data Science

# Gartner's Hype Cycle





1

**BUILD  
A MACHINE LEARNING MODEL  
IN JUST THREE  
QUICK AND EASY STEPS  
USING [...]!!!**

– Most tutorials

# How to Become a Data Scientist?

## How to Draw An Owl



1. Draw some circles



2. Draw the rest of the owl

**50%** of analytic projects fail.

– Gartner, 2015

# Data + Machine Learning = Profit





On September 21, 2009, the grand prize of **US\$1,000,000** was given to the BellKor's Pragmatic Chaos team which bested Netflix's own algorithm for predicting ratings by 10.06%.



“[T]he additional accuracy **gains** that we measured did not seem to justify the engineering **effort** needed to bring them into a production environment.”



Netflix Technology Blog

Learn more about how Netflix designs, builds, and operates our systems and engineering organizations

Apr 5, 2012

## Analytic projects fail because...


...they aren't completed within **budget** or on **schedule**,  
or because they fail to deliver the **features** and **benefits**  
that are optimistically agreed on at their outset.

# How to Avoid Failure?

- 1 Build with Organizational Buy-in
- 2 Build with End In Mind
- 3 Build with a Structured Approach

# How to Avoid Failure?

- 1 Build with Organizational Buy-in
- 2 Build with End In Mind
- 3 Build with a Structured Approach






A conceptual diagram for Data Science. A central white cloud contains the text "Data Science" in bold red font. Radiating from this central cloud are twelve lines, each ending in a white icon. The icons represent various data science concepts: a laptop with a bar chart and magnifying glass, a database cylinder, a cloud with up/down arrows, a document with a bar chart, a server rack, a line graph with a magnifying glass, a magnifying glass over a bar chart, three server units, a pie chart, a database cylinder with a padlock, a line graph with an upward arrow, and a database cylinder with a magnifying glass. The background is a solid reddish-pink color with faint white cloud shapes.

# Data Science


# The Blind Men and the Elephant



It was six men of Indostan  
To learning much inclined,  
Who went to see the Elephant  
(Though all of them were blind),  
That each by observation  
Might satisfy his mind.

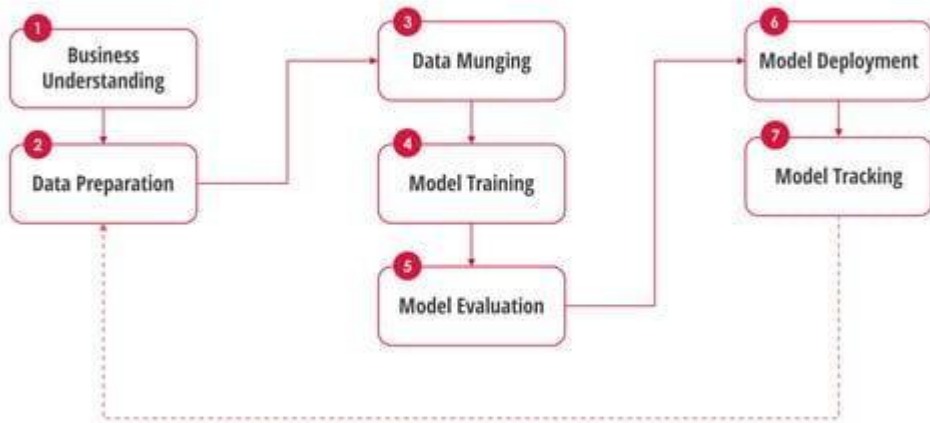
And so these men of Indostan  
Disputed loud and long,  
Each in his own opinion  
Exceeding stiff and strong,  
Though each was partly in the right  
And all were in the wrong!

- John Godfrey Saxe

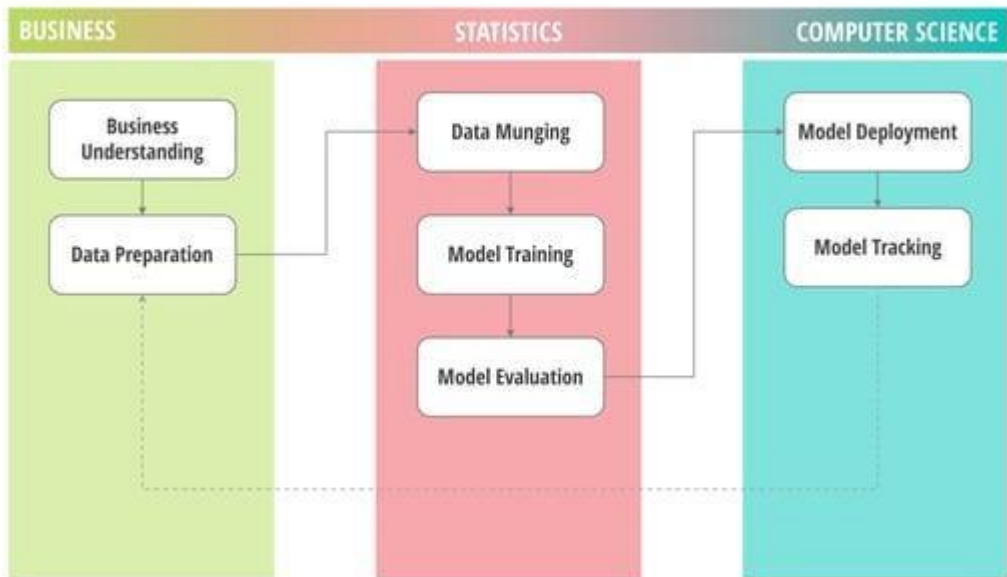




# Data Science Process



# The Data Science Process



Business  
Understanding

Data  
Preparation

Data  
Munging

Model  
Training

Model  
Evaluation

Model  
Deployment

Model  
Tracking

Far better  
an **approximate** answer to the **right** question  
than  
an **exact** answer to the **wrong** question.

– John Tukey

1

DETERMINE

2

UNDERSTAND

3

MAP

## What does the client want to achieve?

1

**DETERMINE**

### Primary Objective

- Reduce attrition
- Customized targeting
- Plan future media spend
- Prevent fraud
- Recommend Products

2

UNDERSTAND

3

MAP

1

DETERMINE

2

UNDERSTAND

- Understand **success criteria**.
  - Specific, measurable, time-bound
- List **assumptions, constraints, and important factors**.
- Identify **secondary or competing objectives**.
- Study **existing solutions** (if any).

3

MAP

1

DETERMINE

2

UNDERSTAND

3

MAP

## Business Objective → Technical Objective

- State the **project objective(s) in technical terms**.
- Describe how the data science project will **help solve the business problem**.
- Explore **successful scenarios**.



1

DETERMINE


2

UNDERSTAND

3

MAP

OBJECTIVE	TECHNIQUE	EXAMPLES
Predict Values	Regression	Linear regression, Bayesian regression, Decision Trees
Predict Categories	Classification	Logistic regression, SVM, Decision Trees
Predict Preference	Recommender System	Collaborative / Content-based filtering
Discover groups	Clustering	<i>k</i> -means, Hierarchical clustering
Identify unusual data points	Anomaly Detection	<i>k</i> -NN, One-class SVM
...		



If all you have is a **hammer**  
then everything looks like a **nail**.




- **Primary Objective:** Prevent attrition → Increase subscription renewals
- **Competing Objective:** High value customers are also targeted for up-sell
- **Constraints:** Avoid targeting customers too close to their contract expiration
- **Success Criteria:** Current renewal rate = 65% → Improve by 8% within the next quarter
- **Existing Solution:** Business-rule-based targeting
- **Data Science Objective:** Build a **binary classification model** to identify customers who are **not likely to renew** their subscriptions at least **three months in advance** of their contract expiration.
- **Success Scenario:** The model correctly identifies **80%** of the future attritors; the promotional campaign targets all likely attritors, and successfully converts **20%** of them into non-attritors.

## Project Plan

- Duration
- Inventory of resources
- Tools and techniques
- Risks and contingencies
- Costs and benefits
- Milestones

**The thought that disaster is impossible often leads to an unthinkable disaster.**

– Gerald Weinberg



Titanic at Southampton docks, prior to departure

1 IDENTIFY

2 COLLECT

3 ASSESS

4 VECTORIZE

## 1 IDENTIFY

- **Data sources, formats**
  - Database, Streaming API's, Logs, Excel files, Websites, etc.
- **Entity Relationship Diagram (ERD)**
- Identify **additional data sources**.
  - Demographics data appends,
  - Geographical data,
  - Census data, etc.
- Identify **relevant data**.
- Record **unavailable data**.
- How long a history is available and one should use?

## 2 COLLECT

## 3 ASSESS

## 4 VECTORIZE

1 IDENTIFY

2 COLLECT

- Access or acquire all relevant data in **a central location**
- **Quality control checks and tests**
  - File formats, delimiters
  - Number of records, columns
  - Primary keys

3 ASSESS

4 VECTORIZE



1 IDENTIFY

2 COLLECT

3 ASSESS

4 VECTORIZE

## First look at the data

- **Get familiar** with the data.
- Study **seasonality**.
  - Monthly/weekly/daily patterns
  - Unexplained gaps or spikes
- Detect **mistakes**.
  - Extreme or outlier values
  - Unusual values
  - Special missing values
- Check **assumptions**.
- Review **distributions**.



**Trust, but verify.**

*Tidy datasets*

Happy families are all alike;

Every unhappy family is unhappy in its own way.

*messy dataset      messy*

- Hadley Wickham



1 IDENTIFY

2 COLLECT

3 ASSESS

4 VECTORIZE

**GOAL:** Create the Analysis Dataset

$$y = \begin{pmatrix} y_1 \\ y_2 \\ y_3 \\ \vdots \\ y_n \end{pmatrix}$$

Outcome  
Target / Labels  
Independent Variable

$$X = \begin{pmatrix} x_{11} & x_{12} & x_{13} & \dots & x_{1j} \\ x_{21} & x_{22} & x_{23} & \dots & x_{2j} \\ x_{31} & x_{32} & x_{33} & \dots & x_{3j} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ x_{n1} & x_{n2} & x_{n3} & \dots & x_{nj} \end{pmatrix}$$


Inputs  
Features / Attributes  
Dependent Variables

# Target Definition

- **Churn = 90 days of consecutive inactivity** (for a pre-paid telecom customer)
- What's **inactivity**?
  - Incoming and outgoing calls
  - Data usage
  - Incoming text
  - Promotional texts
  - Voicemail usage
  - Call forwarding
  - Etc.
- Customers may **change their device** or phone number.
  - Churn at the individual (person) level, or at the device (phone) level?
- Customers may return (become active again) after 90 days of inactivity?
- Prediction window
  - Predict 90 days of consecutive inactivity?
  - Would 10 days of consecutive inactivity suffice?
  - How many customers return after  $x$  days of inactivity?
- Fraud, Involuntary churn
- Etc.

# Accurate but not Precise

## CHART TO HELP DETERMINE RISK OF BEAR ATTACK:





# Modeling Sample

- **Historical trends and seasonality**
  - Are there certain timeframes that should be discarded?
  - The model should be generalizable.
- **Eligible, relevant population**
  - Must align with the business goals
- **Eligible, relevant markets**
  - Must align with the business goals
  - E.g., within a certain drive-time distance
- **Outdated products or events**



# Selection Bias



Abraham Wald's Work on Aircraft Survivability  
*Journal of the American Statistical Association* Vol. 79, No. 386 (Jun., 1984)

# Information Leakage



- The leading indicators must be calculated from the timeframe *leading up to* the event – it must not overlap with the prediction window.
- Beware of proxy events, e.g., future bookings.



- **Attribute creation**
  - Derived attributes: Household income / Number of adults = Income per adult
- **Brainstorm with team members** (both technical and non-technical)

$$X = \begin{pmatrix} x_{11} & x_{12} & x_{13} & \dots & x_{1j} \\ x_{21} & x_{22} & x_{23} & \dots & x_{2j} \\ x_{31} & x_{32} & x_{33} & \dots & x_{3j} \\ \vdots & \vdots & \vdots & & \vdots \\ \vdots & \vdots & \vdots & & \vdots \\ \vdots & \vdots & \vdots & & \vdots \\ x_{n1} & x_{n2} & x_{n3} & \dots & x_{nj} \end{pmatrix}$$

# Data Aggregation

CUSTOMER ID	PURCHASE DATE
1001	02-12-2015:05:20:39
1001	05-13-2015:12:18:09
1001	12-20-2016:00:15:59
1002	01-19-2014:04:28:54
1003	01-12-2015:09:20:36
1003	05-31-2015:10:10:02
...	...



CUSTOMER ID	$x_1$	$x_2$	...	$x_j$
1001	...	...		...
1002	...	...		...
1003	...	...		...
...	...	...	...	...

1. Number of transactions (Frequency)
2. Days since the last transaction (Recency)
3. Days since the earliest transaction (Tenure)
4. Avg. days between transaction
5. # of transactions during weekends
6. % of transactions during weekends
7. # of transactions by day-part (breakfast, lunch, etc.)
8. % of transactions by day-part
9. Days since last transaction / Avg. days between transactions
- 10....

1 IDENTIFY

2 COLLECT

3 ASSESS

4 VECTORIZE

**OUTPUT:** The Analysis Dataset

$$y = \begin{pmatrix} y_1 \\ y_2 \\ y_3 \\ \vdots \\ y_n \end{pmatrix} \quad X = \begin{pmatrix} x_{11} & x_{12} & x_{13} & \dots & x_{1j} \\ x_{21} & x_{22} & x_{23} & \dots & x_{2j} \\ x_{31} & x_{32} & x_{33} & \dots & x_{3j} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ x_{n1} & x_{n2} & x_{n3} & \dots & x_{nj} \end{pmatrix}$$

kaggle



Business  
Understanding

Data  
Preparation

Data  
Munging

Model  
Training

Model  
Evaluation

Model  
Deployment

Model  
Tracking

Time  
Spent

80%

20%

Data  
Munging

Model  
Building

Give me **six** hours to chop down a tree  
and I will spend the first **four** sharpening the axe.

– Anonymous

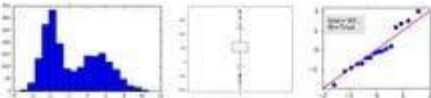
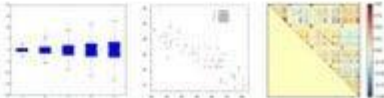
- **Descriptive statistics**
  - Review with the client
- **Correlation analysis**
  - Review with the client
  - Watch out for data leakage
- **Impute missing values**
- **Trim extreme values**
- **Process categorical attributes**
- **Transformations** (square, log, etc.)
  - Binning / variable smoothing
- **Multicollinearity**
  - Reduce redundancy
- **Create additional feature**
- **Interactions**
- **Normalization** (scaling)



via @boykis

**Machine learning experts display  
cleaned data samples  
in preparation for modeling.**

Annibale Carracci, c. 1600

	Univariate	Multivariate
Non-Graphical	<ul style="list-style-type: none"> <li>○ Categorical: Tabulated frequencies</li> <li>○ Quantitative:               <ul style="list-style-type: none"> <li>○ Central tendency: mean, median, mode</li> <li>○ Spread: Standard deviation, inter-quartile range</li> <li>○ Skewness and kurtosis</li> </ul> </li> </ul>	<ul style="list-style-type: none"> <li>○ Cross-tabulation</li> <li>○ Univariate statistics by category</li> <li>○ Correlation matrices</li> </ul>
Graphical	<ul style="list-style-type: none"> <li>○ Histograms</li> <li>○ Box plots, stem-and-leaf plots</li> <li>○ Quantile-normal plots</li> </ul> 	<ul style="list-style-type: none"> <li>○ Univariate graphs by category (e.g., side-by-side box-plots)</li> <li>○ Scatterplots</li> <li>○ Correlation matrix plots</li> </ul> 

- **Feature Reduction:** The process of selecting a subset of features for use in model construction
  - Useful for both supervised and unsupervised learning problems

**Art is the elimination of the unnecessary.**

– Pablo Picasso



## Feature Reduction: Why

- **True dimensionality <<< Observed dimensionality**
  - The abundance of redundant and irrelevant features
- **Curse of dimensionality**
  - With a fixed number of training samples, the predictive power reduces as the dimensionality increases. [Hughes phenomenon]
  - With  $d$  binary variables, the number of possible combinations is  $O(2^d)$ .
- **Goal of the Analysis**
  - Descriptive → Diagnostic → Predictive → Prescriptive

Hindsight	Insight	Foresight
-----------	---------	-----------
- **Law of Parsimony [Occam's Razor]**
  - Other things being equal, simpler explanations are generally better than complex ones.
- **Overfitting**
- **Execution time (Algorithm and data)**

## Feature Reduction Techniques



A practical guide to dimensionality reduction techniques – Vishal Patel


1. Percent missing values
2. Amount of variation
3. Pairwise correlation
4. Multicollinearity
5. Principal Component Analysis (PCA)
6. Cluster analysis
7. Correlation (with the target)
8. Forward selection
9. Backward elimination
10. Stepwise selection
11. LASSO
12. Tree-based selection

- Try **more than one** machine learning technique.
- Fine-tune **parameters**.
- Assess **model performance**.
- Avoid **Over-fitting**.




# Assess Model Performance

**Type I error**  
(false positive)



**Type II error**  
(false negative)



- **New Age:** Area Under the ROC Curve (AUC), Confusion Matrix, Precision, Recall, Log-loss, etc.
- **Old School:** Model Lift, Model Gains, Kolmogorov-Smirnov (KS), etc.

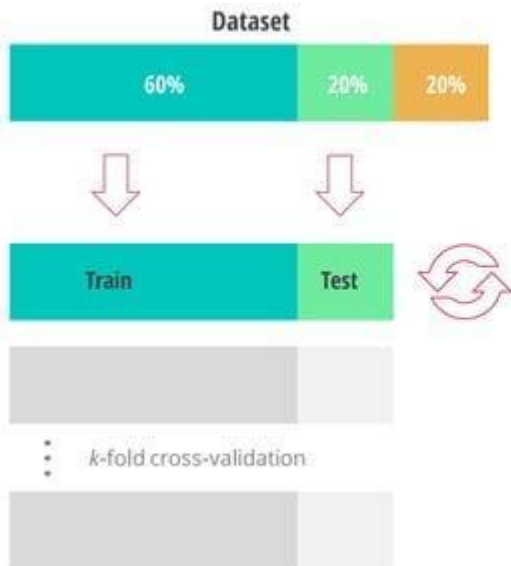
**When a measure becomes a target,  
it ceases to be a good measure.**

Goodhart's law





Photo Courtesy: @bucaris

# Tri-fold Partition



- Fine-tune and select the best model based on **Train** + **Test** sets.
- Evaluate the chosen algorithm on the **Validation** set (i.e., completely unseen data).

# Bias-Variance Tradeoff



**With four parameters I can fit an elephant,  
and with five I can make him wiggle his trunk.**

- John von Neumann



1

MODEL SELECTION

2

ASSESSMENT

3

PRESENTATION

## 1 MODEL SELECTION

- Law of Parsimony (Occam's Razor)
- Model execution time
- Deployment complexity

## 2 ASSESSMENT

## 3 PRESENTATION

Build the simplest solution that can adequately answer the question.

1 MODEL SELECTION

2 ASSESSMENT

Dataset

20%

Validation

Temporal  
or  
Random

3 PRESENTATION

## 1 MODEL SELECTION

## 2 ASSESSMENT

## 3 PRESENTATION

- AUC, etc.
- Cumulative Gains Chart / Lift Chart
- Compare against existing business rules/model
- Predictor Importance
- Each predictor's relationship with the target
- Reason-coding
- Model usage recommendations
  - Decile reports
- Personify
- Model peer-review (Quality Control)

Interpret results as they relate to the business application.

- **Model production cycle**
- **Scoring code, or publish model as a web service**
  - Hand-off
- **Model Documentation** (Technical Specifications)
  - Data preparation, transformations, imputations, parameter settings, etc.
- **Reproducibility**
  - Docker containers
- **Model Persistence vs. Model Transience**

# Model Persistence vs. Model Transience



1 MONITOR

2 MAINTAIN

3 TEST

1

**MONITOR**

- **Model decay tracking (monitoring) plan**

- Model performance over time
- Predictor distribution

2

**MAINTAIN**

3

**TEST**



1 MONITOR

2 **MAINTAIN**

- Model maintenance plan
- Adding new data sources
- Version control

3 TEST

1 MONITOR

2 MAINTAIN

3 TEST

- Campaign Set-up and Execution
  - Experimental Design (A/B tests, Fractional Factorial)

# Experimental Design

	Marketing Treatment	No Treatment
Selection Based on Model	<b>A</b> Test	<b>B</b> Selection Hold-out
No Selection (Random)	<b>C</b> Control	<b>D</b> Random Hold-out



# Data Science Process: Recap

Business Understanding	Data Preparation	Data Munging	Model Training	Model Evaluation	Model Deployment	Model Tracking
Determine	Identify	Impute	Train	Evaluate	Deploy	Monitor
Understand	Collect	Transform	Assess	Peer Review	Document	Maintain
Map	Assess	Reduce	Select	Present		Test
	Vectorize					

DISCUSS

COLLATE

WRANGLE

PERFORM

COMMUNICATE

EXECUTE

TRACK

# Process as Proxy

“Good process **serves you** so you can serve customers.

But if you're not watchful, the process can become the **proxy** for the result you want.

You stop looking at outcomes and just make sure you're doing the process right.

Gulp.

It's always worth asking, **do we own the process** or **does the process own us?**”

– Jeff Bezos

# THANK YOU!

vishal@derive.io

[www.linkedin.com/in/VishalJP](https://www.linkedin.com/in/VishalJP)

**DΣRIVE**