

Trade-offs in Maintaining Consistency

- Maintaining consistency should balance between the strictness of consistency versus efficiency
 - How much consistency is “good-enough” depends on the application

Loose Consistency



Strict Consistency



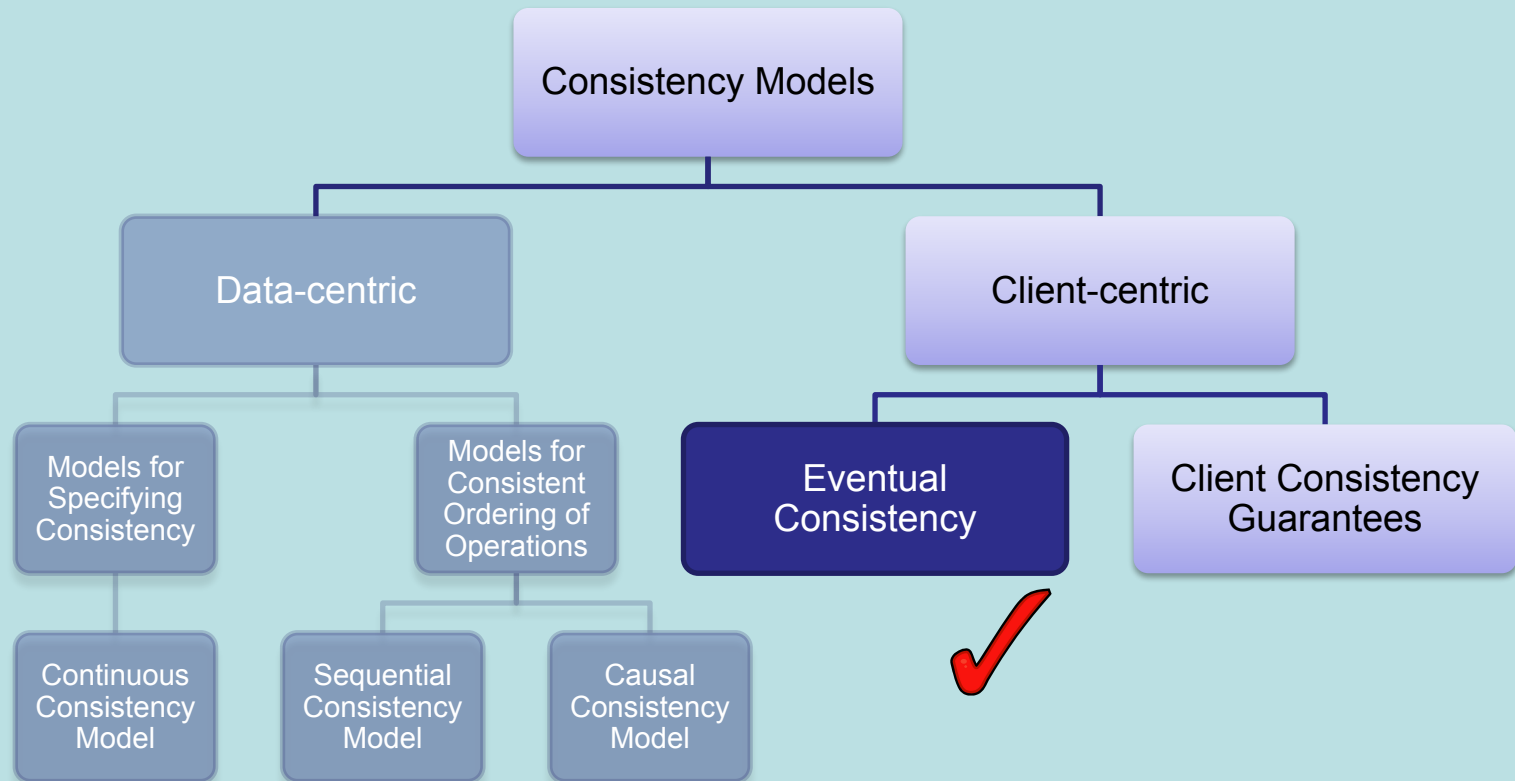
Easier to implement,
and is efficient

Generally hard to implement,
and is inefficient

Client-Centric Consistency Models

- Data-centric models lead to excessive overheads in applications where:
 - a majority operations are reads, and
 - updates occur frequently, and are often from one client process
- For such applications, a weaker form of consistency called *Client-centric Consistency* is employed for improving efficiency
- Client-centric consistency models specify two requirements:
 1. **Eventual Consistency**
 - All the replicas should *eventually* converge on a final value
 2. **Client Consistency Guarantees**
 - Each client processes should be guaranteed some level of consistency while accessing the data value from different replicas

Overview



Eventual Consistency

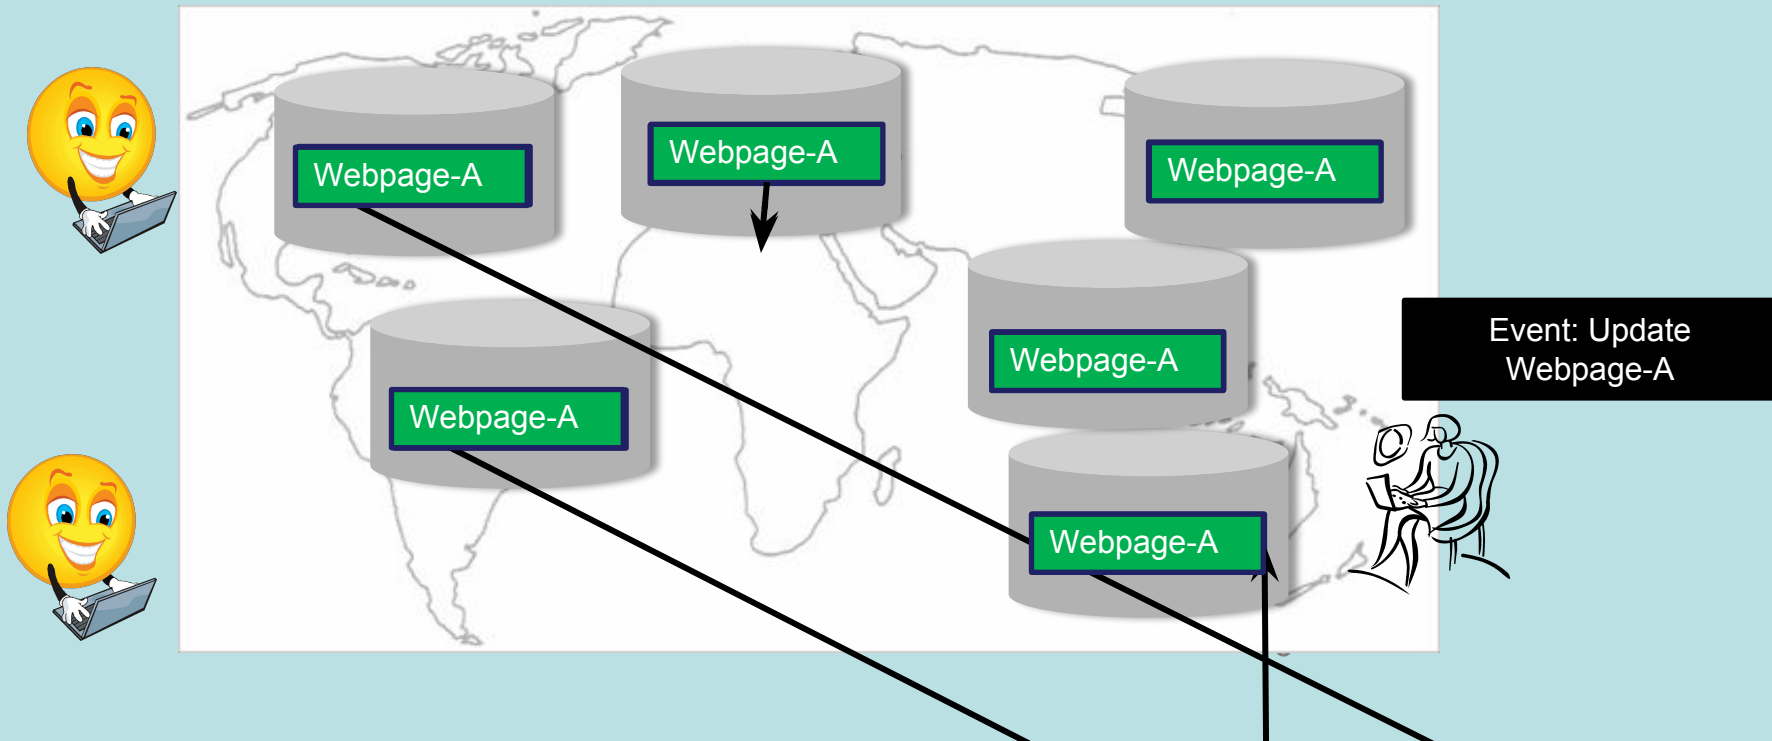
- Many applications can tolerate a inconsistency for a long time
 - Webpage updates, Web Search – Crawling, indexing and ranking, Updates to DNS Server
- In such applications, it is acceptable and efficient if replicas in the data-store rarely exchange updates
- A data-store is termed as *Eventually Consistent* if:
 - All replicas will gradually become consistent in the absence of updates
- Typically, updates are propagated infrequently in eventually consistent data-stores

Designing Eventual Consistency

- In eventually consistent data-stores,
 - *Write-write conflicts* are rare
 - Two processes that write the same value are rare
 - Generally, one client updates the data value
 - e.g., One DNS server updates the name to IP mapping
 - Such rare conflicts can be handled through simple mechanisms, such as mutual exclusion
 - *Read-write conflict* are more frequent
 - Conflicts where one process is reading a value, while another process is writing a value to the same variable
 - Eventual Consistency Design has to focus on efficiently resolving such conflicts


Challenges in Eventual Consistency

- Eventual Consistency is not good-enough when the client process accesses data from different replicas
 - We need consistency guarantees for a single client while accessing the data-store



Client Consistency Models

- Client-centric consistency provides guarantees for a single client for its accesses to a data-store
- Example: Providing consistency guarantee to a client process for data x replicated on two replicas. Let x_i be the local copy of a data x at replica I_i .



$WS(x_1)$ = Write Set for x_1 = Series of ops being done at some replica that reflects how I_1 updated x_1 till this time

$WS(x_1; x_2)$ = Write Set for x_1 and x_2 = Series of ops being done at some replica that reflects how I_1 updated x_1 and, later on, how x_2 is updated on I_2



= Replica i

$R(x_i) b$

= Read variable x at
replica i ; Result is b

$W(x) b$

= Write variable x at
replica i ; Result is b

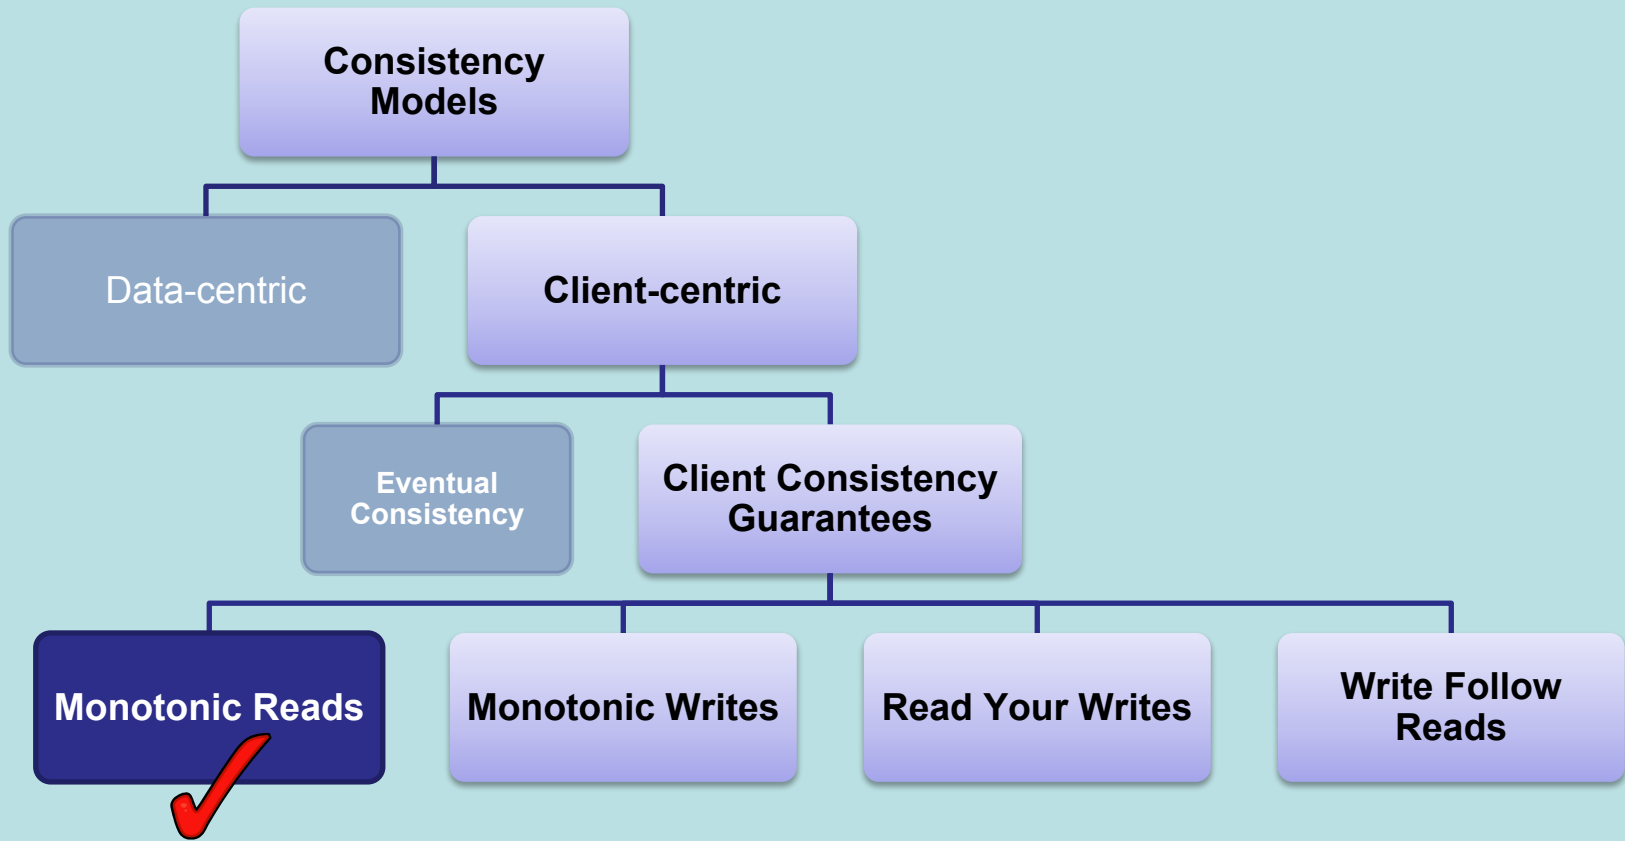
$WS(x_i)$

= Write Set

Client Consistency Models

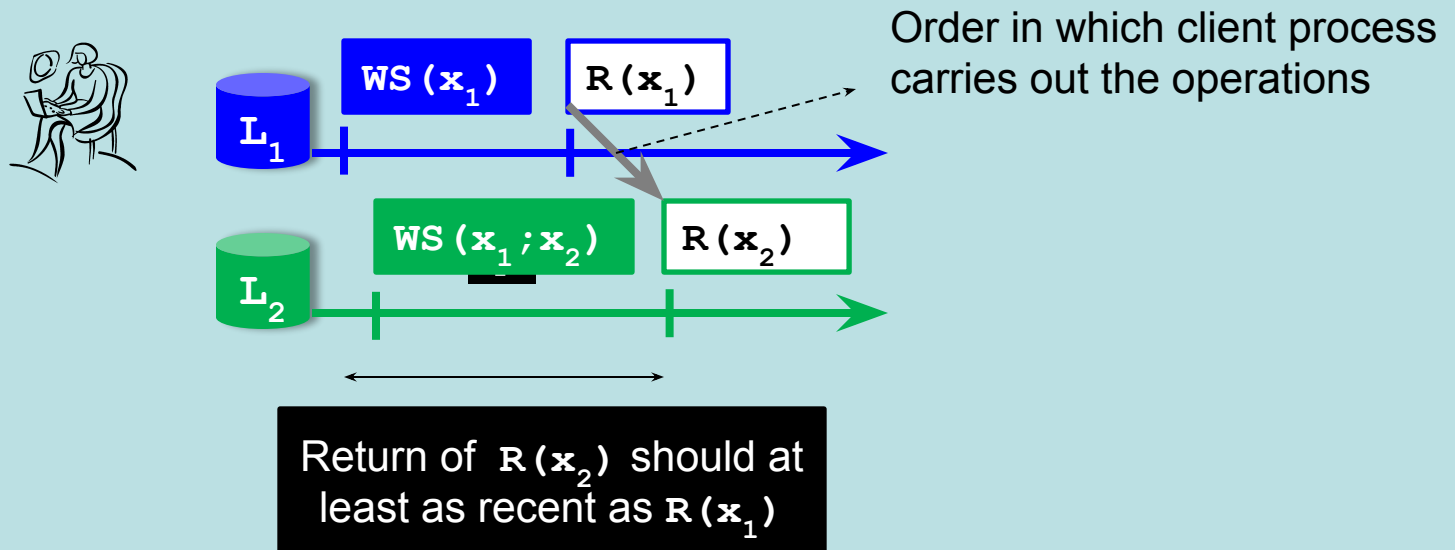
- We will study four types of client-centric consistency models¹
 1. Monotonic Reads
 2. Monotonic Writes
 3. Read Your Writes
 4. Write Follow Reads

Overview



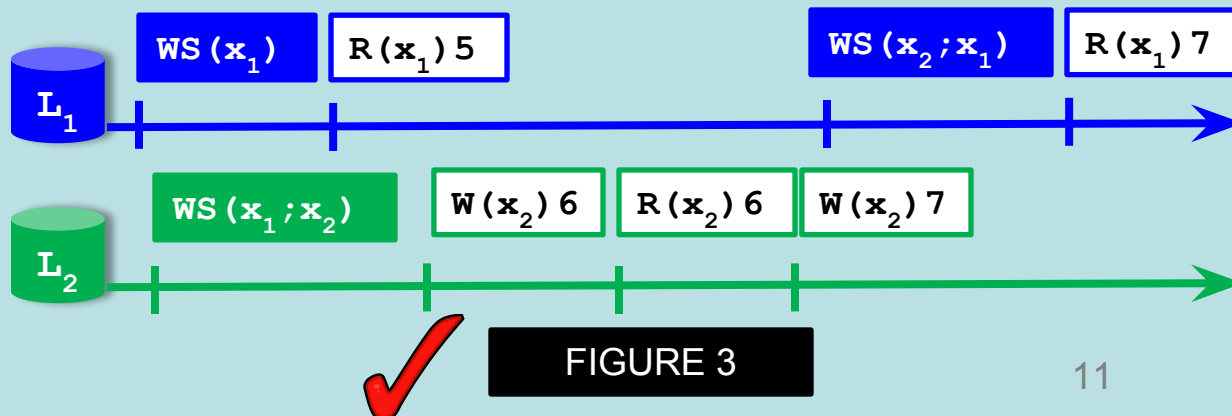


Monotonic Reads

- The model provides guarantees on successive reads
- If a client process reads the value of data item x , then any successive read operation by that process should return the same or a more recent value for x

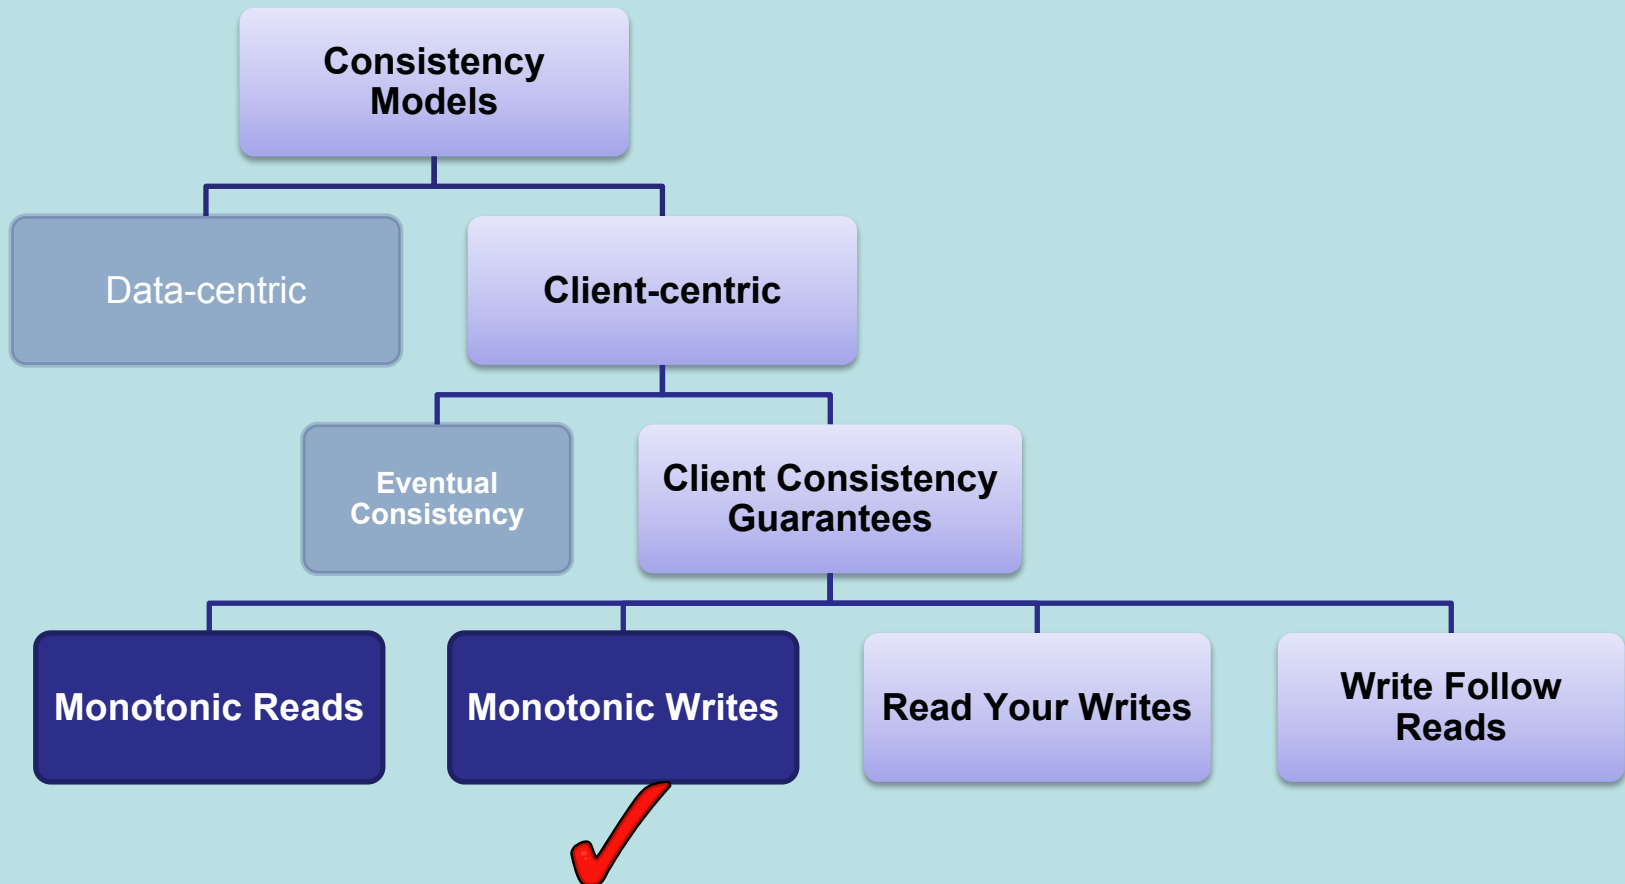


Monotonic Reads – Puzzle

Recognize data-stores that provide monotonic read guarantees




Overview




Monotonic Writes

- This consistency model assures that writes are monotonic
- A write operation by a client process on a data item x is completed before any successive write operation on x by the same process
 - A new write on a replica should wait for all old writes on any replica



$W(x_2)$ operation should be performed only after the result of $W(x_1)$ has been updated at I_2

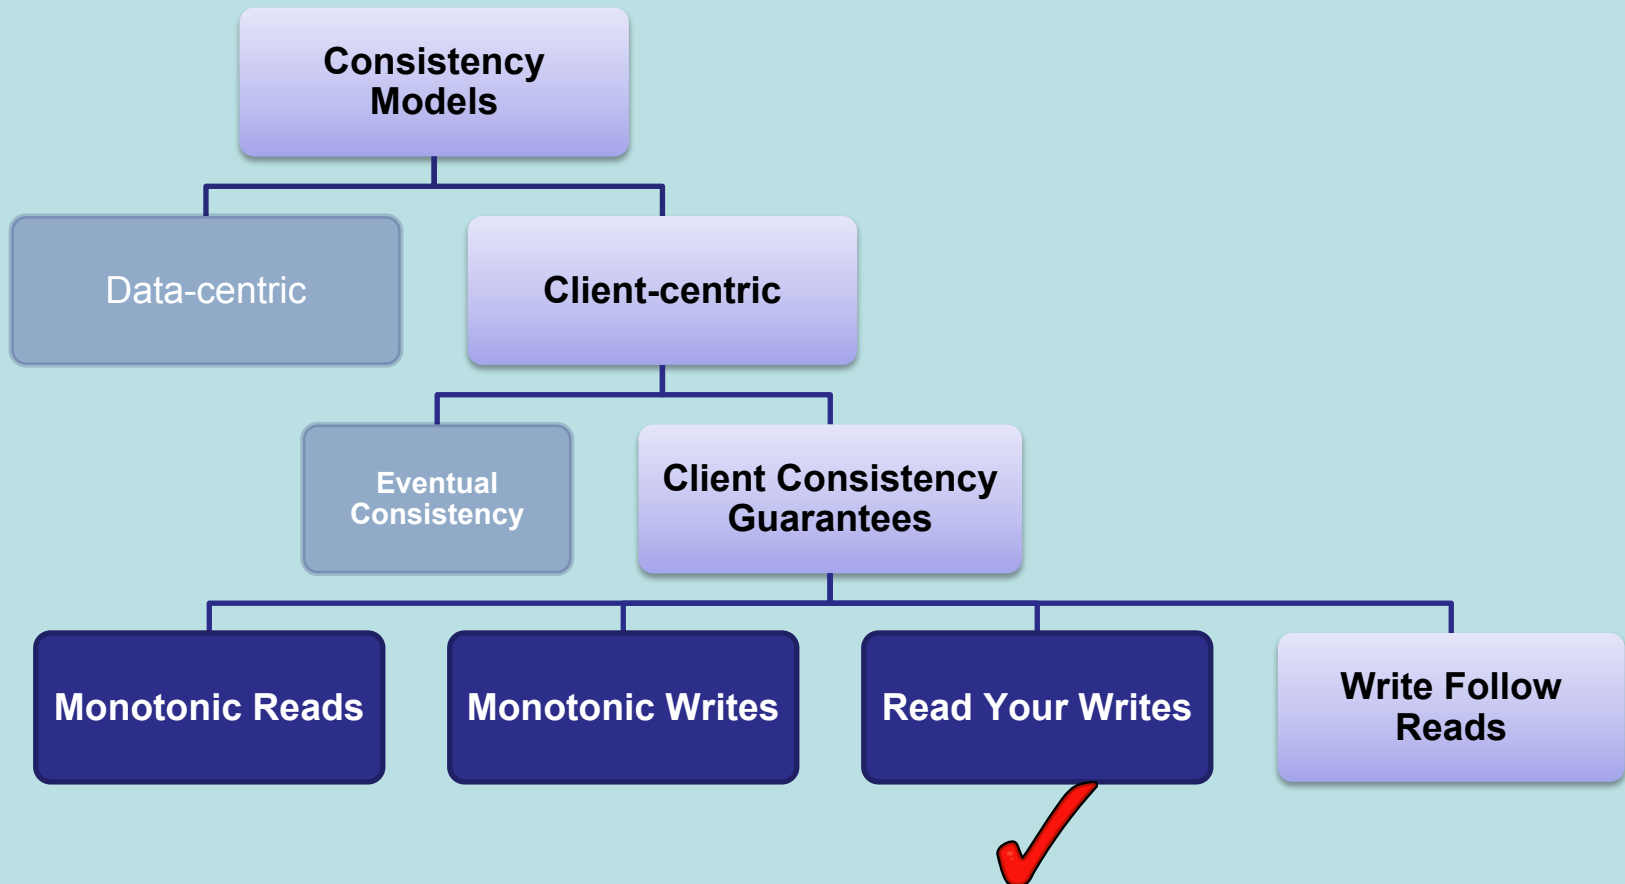


The data-store does not provide monotonic write consistency

Monotonic Writes – An Example

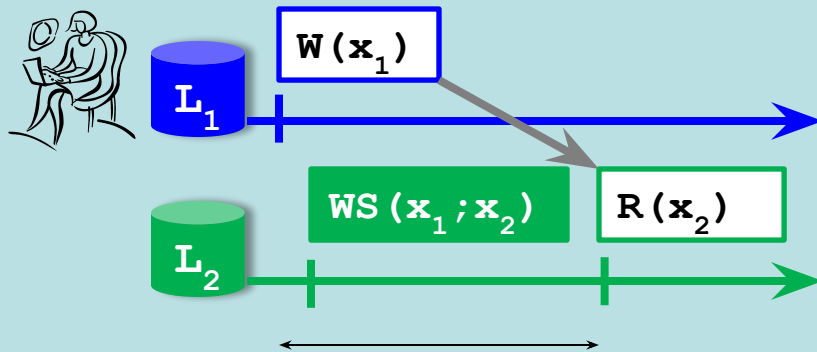
- Example: Updating individual libraries in a large software source code which is replicated
 - Updates can be propagated in a lazy fashion
 - Updates are performed on a part of the data item
 - Some functions in an individual library is often modified and updated
 - Monotonic writes: If an update is performed on a library, then all preceding updates on the same library are first updated
- Question: If the update overwrites the complete software source code, is it necessary to update all the previous updates?

Overview




Read Your Writes

- The effect of a write operation on a data item x by a process will always be seen by a successive read operation on x by the same process
- Example scenario:
 - In systems where password is stored in a replicated data-base, the password change should be seen immediately

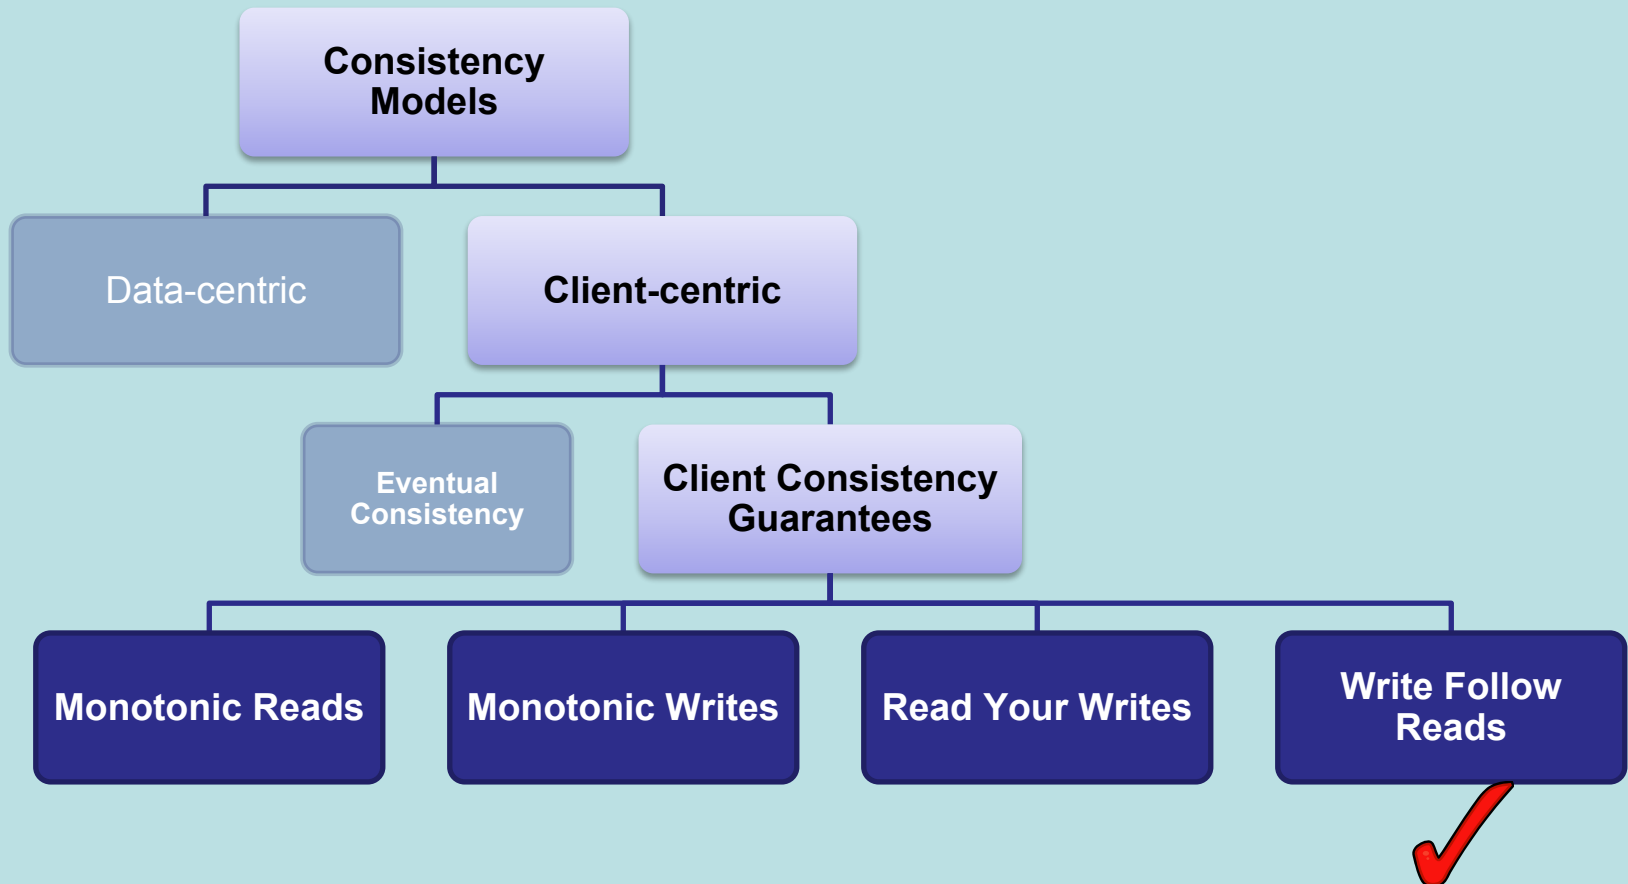


$R(x_2)$ operation should be performed only after the updating the Write Set $WS(x_1)$ at I_2



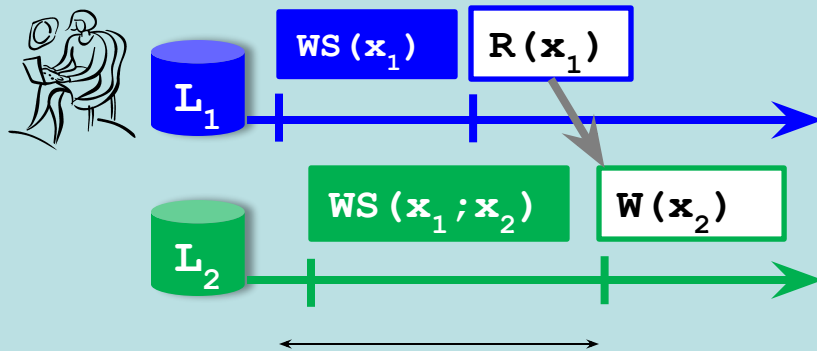
A data-store that does not provide *Read Your Write* consistency

Overview




Write Follow Reads

- A write operation by a process on a data item x following a previous read operation on x by the same process is guaranteed to take place on the same or a more recent value of x that was read
- Example scenario:
 - Users of a newsgroup should post their comments only after they have read all previous comments



$W(x_2)$ operation should be performed only after the all previous writes have been seen



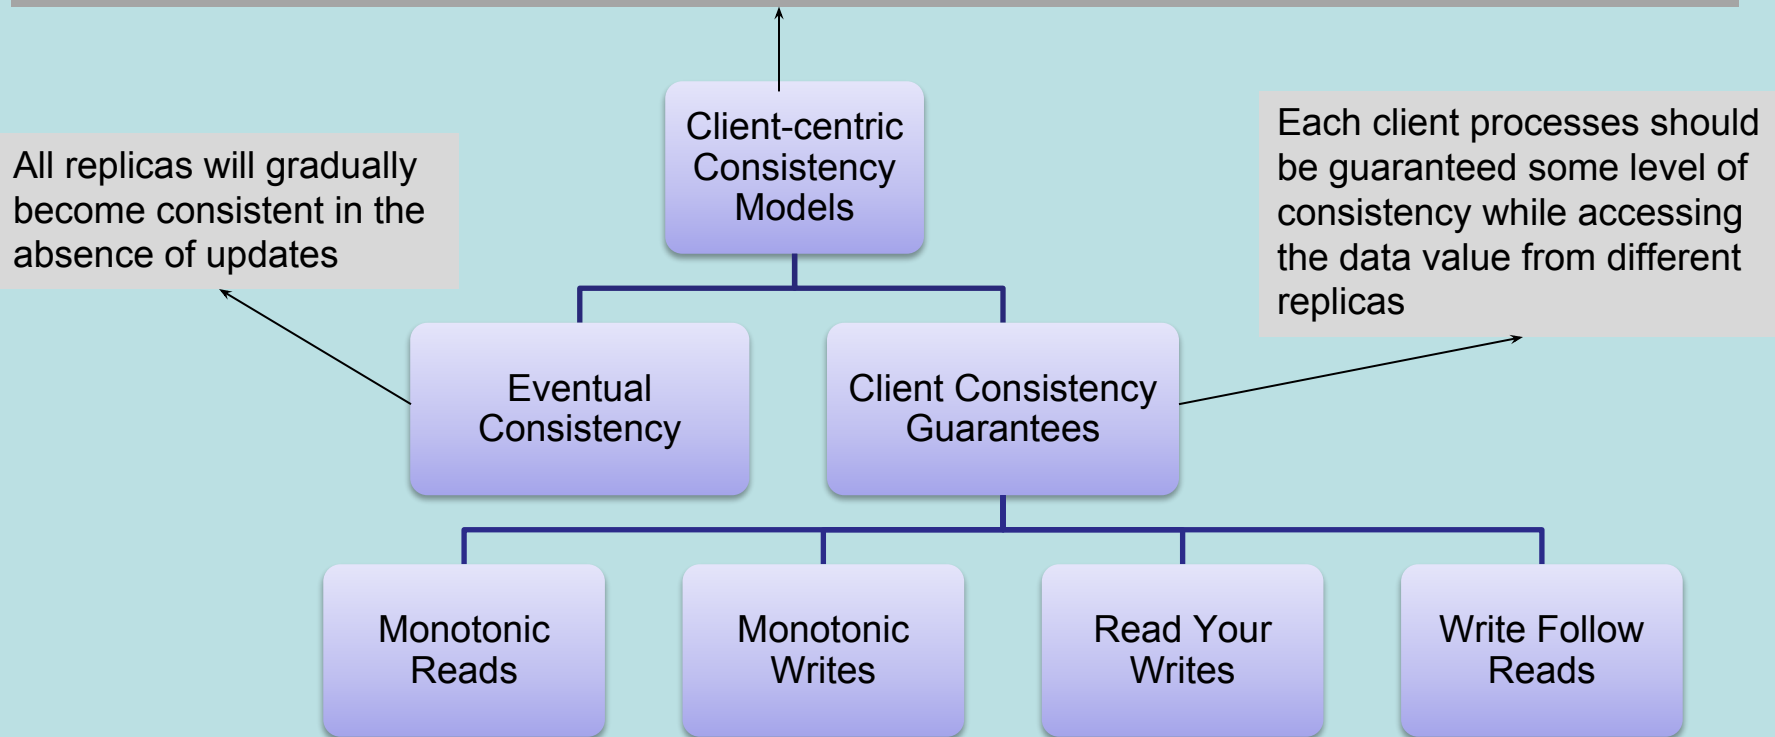
A data-store that does not guarantee Write Follow Read Consistency Model

Summary of Client-centric Consistency Models

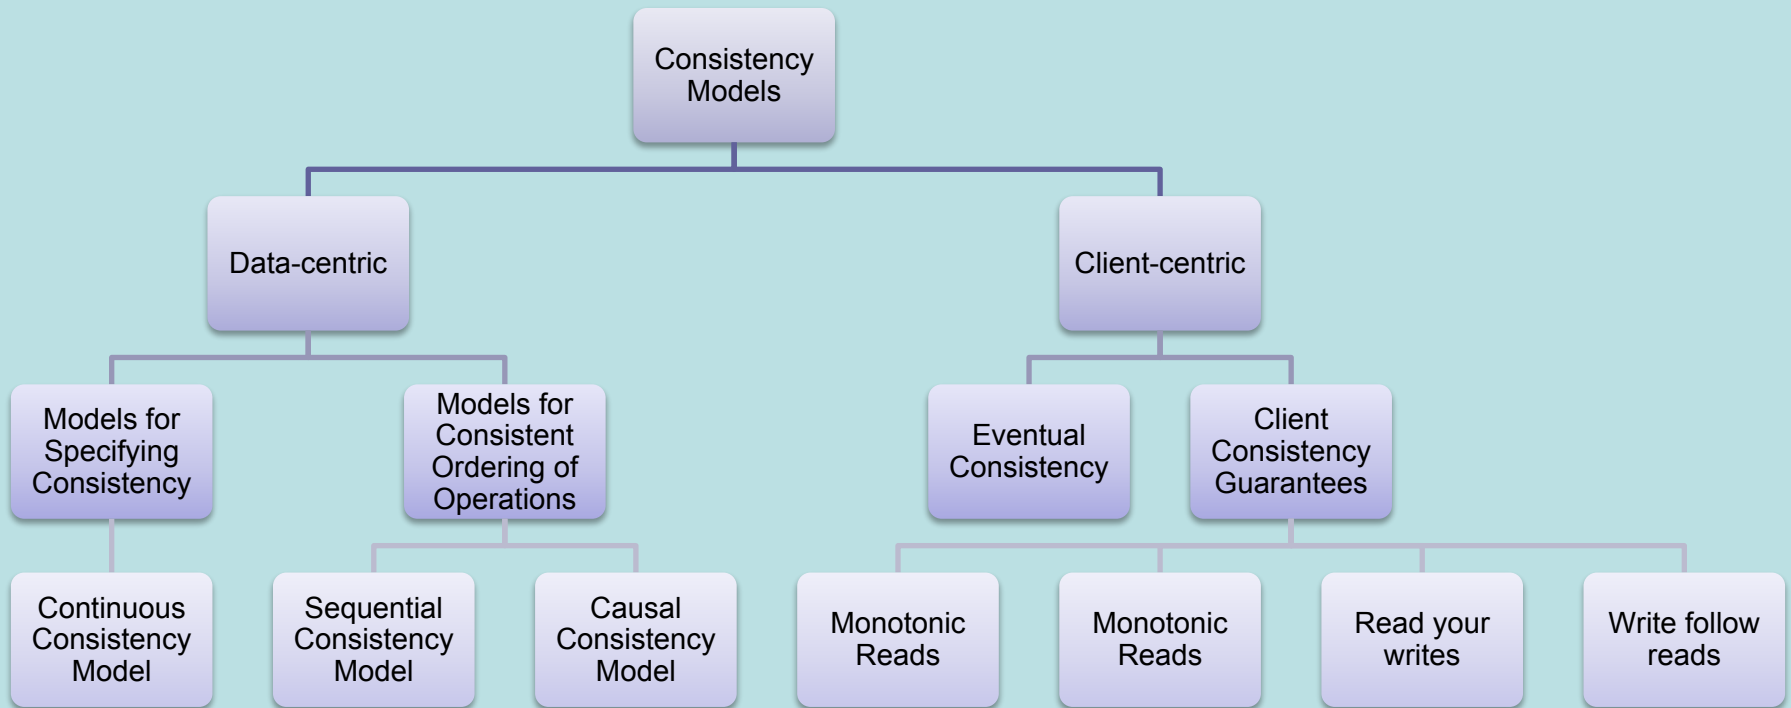
Client-centric Consistency Model defines how a data-store presents the data value to an individual client when the client process accesses the data value across different replicas.

It is generally useful in applications where:

- one client always updates the data-store.
- read-to-write ratio is high



Topics covered in Consistency Models



Summary of Consistency Models

- Different applications require different levels of consistency
 - Data-centric consistency models
 - Define how replicas in a data-store maintain consistency
 - Client-centric consistency models
 - Provide an efficient, but weaker form of consistency when
 - Here, one client process updates the data item, and many processes read the replica

Next Class

- Replica Management
 - Describes where, when and by whom replicas should be placed
- Consistency Protocols
 - We study “how” consistency is ensured in distributed systems